# BUILDING AN EFFICIENT ETL PIPELINE

By:

Janani III/CSE Allen III/CSE Poonam III/IT Akaash III/CSE











## **/TABLE OF CONTENTS**

/01	AIM	/02	WHY THIS PROBLEM STATEMENT
/03	SOLUTION - PHASE 2 OVER PHASE 1	/04	TECHNICAL EXPLANATION
/05	OPTIMIZATION & INNOVATION	/06	VIABILITY
/07	MARKET COMPARISON & ESTIMATION	/08	FUTURE PLANS







To build an efficient **ETL Pipeline** which Extracts, Transforms and Loads data from a destination to another. It must consist of a web dashboard which performs **visualization** on data and helps **reporting errors** occurring in the pipeline.









## **/WHY WE CHOSE THIS PS**

- Interesting problem statement
- ETL, Cloud, Data Warehousing is a new concept to us - we wanted to take it up as a challenge.
- There are not many resources like other problem statements- so we learnt topics from scratch by doing a lot of research.

## INITIAL SOLUTION - what we tried performing in the beginning weeks

- Extracting data from SQL Server and loading data into Postgres
- Automating the pipeline using Airflow (using TaskFlow API) and using DAG
- Performing Incremental Data loading using Destination Change Comparison technique
- Alteryx- data visualization tool that works excellently with ETL

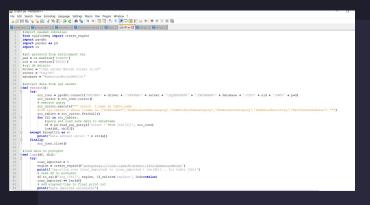






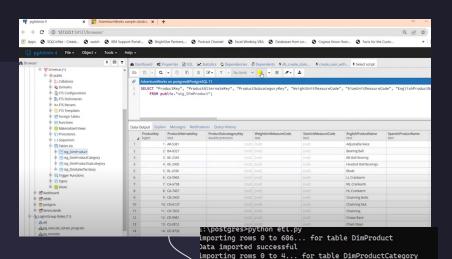
#### **INITIAL SOLUTION 1**

1. Making a basic ETL pipeline using Python



2. Results were produced in Postgre from SSMS

⇒ETL successful



Data imported successful

Data imported successful

Data imported successful

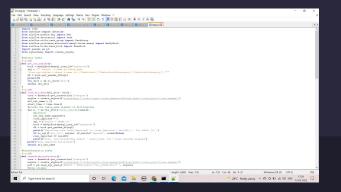
Data imported successful

Importing rows 0 to 37... for table DimProductSubcategory

importing rows 0 to 60398... for table FactInternetSales

importing rows 0 to 11... for table DimSalesTerritory

3. DAG code for automation purposes











## **INITIAL SOLUTION 1 PROBLEMS**

#### Postgres

- Loading takes a long time for big data
- Multiple data access becomes difficult
- Low reading speed

#### Airflow

Airflow doesn't have versioning- if you delete a task from DAG code and redeploy it,
 metadata related to task will be lost









## **INITIAL SOLUTION 2**

#### PROBLEMS WRT DOCKER

- Writing ETL pipeline in Python.
- Apply functional programming in Data Engineering.
- Use a proper object oriented code design and meta file for job control.
- Implement a pipeline in Python extracting data
   Using docker image with application code on
   Docker Hub, using Kubernetes and Argo Workflows

- Is best for building microservices for applications
- Not good if software runs in different environments or are from multiple sources
- Lesser speed
- Data storage is complicated
- Graphical applications don't work well
- In Docker files are created inside a container. It is difficult to retrieve out of the container for different kind of processes









## SOUR SOLUTION?

During our last few weeks?

>Using Cloud of course! <</pre>











## **/WHY CLOUD?**

Simplified data process for users	Allows organizations to combine data from multiple sources	Easy optimization			
Handles huge volumes of data and data migration	Easy for visualization and debugging and reduce MTTR	Easy to compute data			
No need of much security patches	Maintenance	Easy to analyze unstructured data			









## **TOOLS WE ARE USING**









#### **DBEAVER**

To create and manage databases across a wide range of database management systems.

#### **REDSHIFT**

Large scale data storage and analysis, and is frequently used to perform large database migrations.

#### **DATA PIPELINE**

helps you reliably process and move data between different AWS compute and storage services, as well as on-premises data sources, at specified intervals.

#### **GLUE**

Uses ETL jobs to extract data from a combination of other Amazon Web Services and incorporates it into data lakes and data warehouses.









## **TOOLS WE ARE USING**









#### **CLOUDWATCH**

An AWS monitoring service you can use to monitor your applications, services, and resources.

#### **QUICKSIGHT**

Use to deliver easy-to-understand insights to the people who you work with, wherever they are. Amazon QuickSight connects to your data in the cloud and combines data from many different sources.

#### **PYSPARK**

A standard ETL tool like PySpark, supports all basic data transformation features like sorting, mapping, joins, operations, etc.

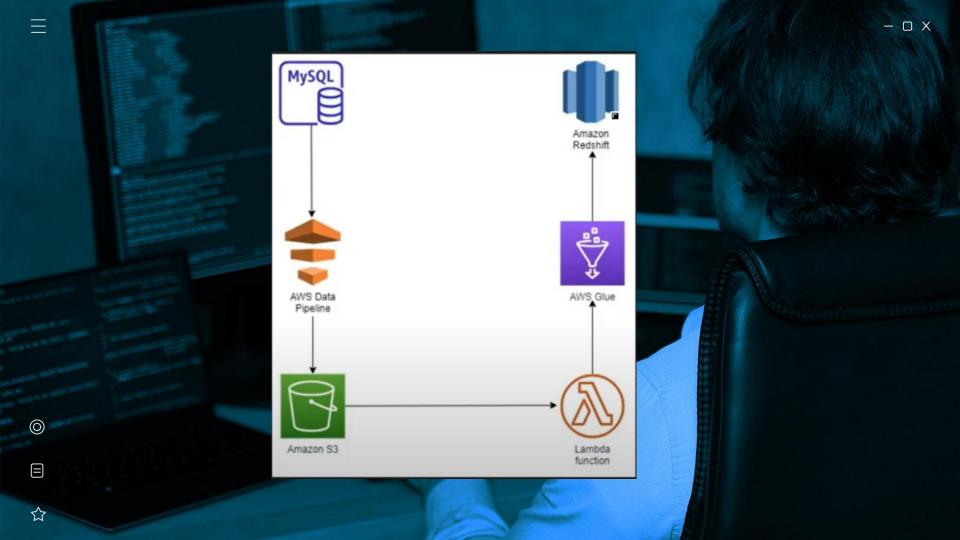
#### **ATHENA**

Performs interactive queries in the web-based cloud storage service, Amazon Simple Storage Service (S3). Athena is used with large-scale data sets. Amazon S3 is designed for online backup and archiving of data











# /TECHNICAL EXPLANATION

OUR SOLUTION EXPLANATION











### TECHNICAL EXPLANATION

- We first made a database connection between DBeaver and RDS console. DBeaver will have the entire database.
- Then we set up Redshift cluster by creating roles, editing IAM roles and VPC and configured it as per the requirement to build the ETL
- Created a S3 bucket which has an initial load and incremental load
- In Data Pipeline we configured DataNodes, table structure and created an EC2 instance



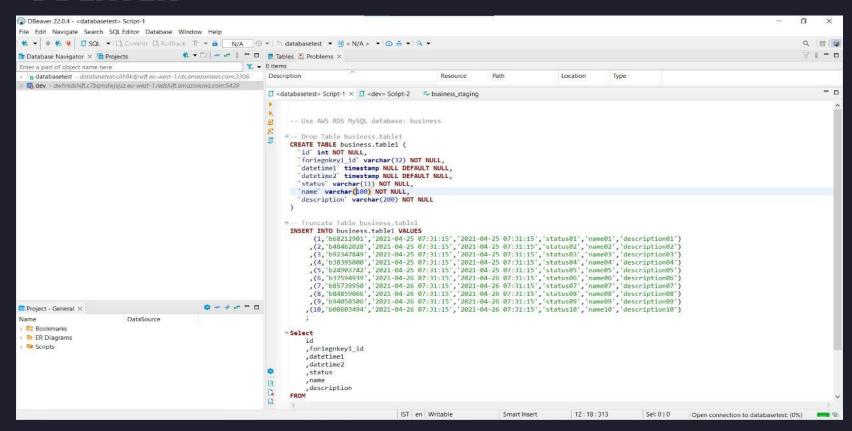






#### $-\Box X$

### **DBEAVER**

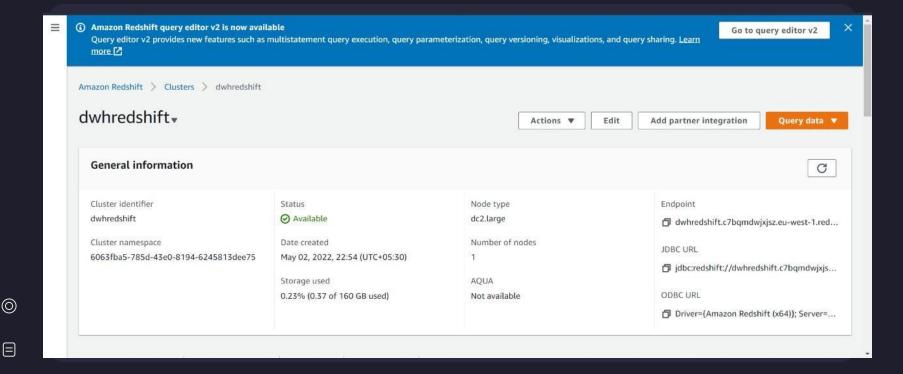








## REDSHIFT CLUSTER







## **/COMPARISON**

	/BIGQUERY	/REDSHIFT
Service	Serverless/Managed	Not Managed
What are they	fully managed cloud-based data warehouse which is designed for handling large-scale data set storage.	fully managed and serverless data warehouse.
Regions Available	Americas – 4 Europe – 3 Asia Pacific – 7	Americas - 8 Europe - 5 Asia Pacific - 9
Storage format	Columnar & uncompressed storage	Columnar & compressed storage
Billing	Not Predictable	Predictable









## **EXPLANATION**



- AWS Redshift connection was established with DBeaver
- Initial data load was performed into Redshift for data load for a base object
- Incremental data load was performed into Redshift for subsequent load process requirements
- Incremental Data Pipeline from MySQL to S3 was performed by altering DataNodes and EC2 Resource
- AWS Glue Job from S3 was then configured with Redshift

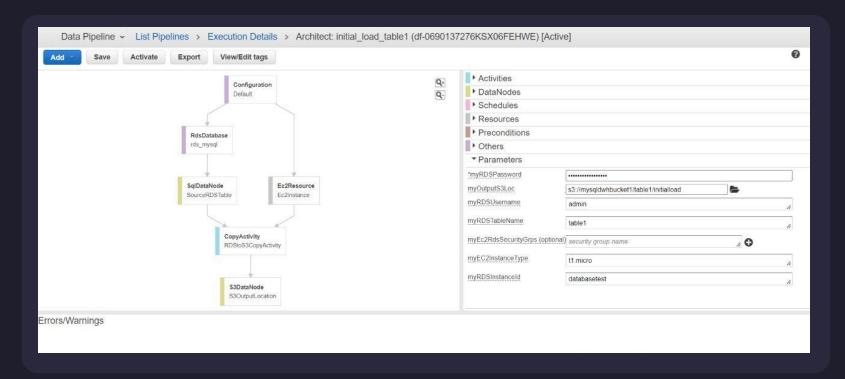






## **DATA PIPELINE**













## **EXPLANATION**

- Lambda Function for AWS Glue Job from S3 was configured to AWS Redshift
- Redshift Snapshot and Restore process performed
- Our ETL was ready
- We then used various tools to visualize and debug our ETL

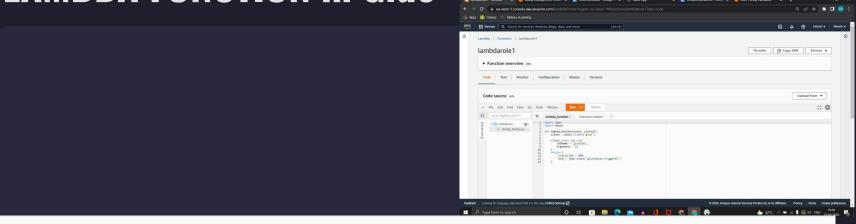








## LAMBDA FUNCTION in Glue





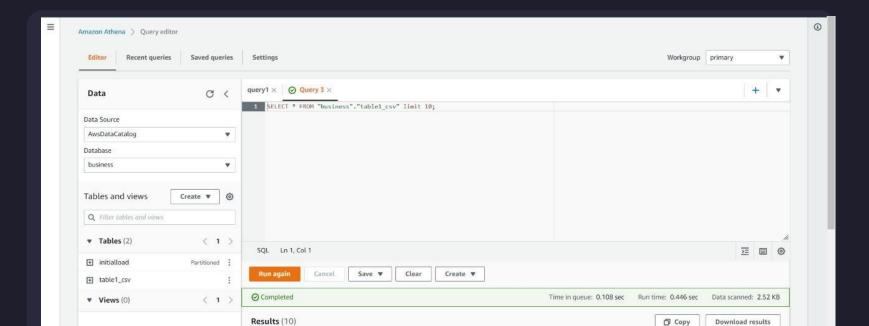






(1) 0

## **ATHENA**



Q Search rows









## **/QUICKSIGHT**



#### /VISUALIZATION

- Interactive dashboards
- Combine a variety of data into a single analysis.and publish dashboards
- Less expensive than Tableau
- Any data can be used



#### /SPICE

- Super-fast parallel in-memory calculation engine
- It's faster and easier to retrieve than having to query the original data source each time
- Rapidly perform advanced calculations and serve the data



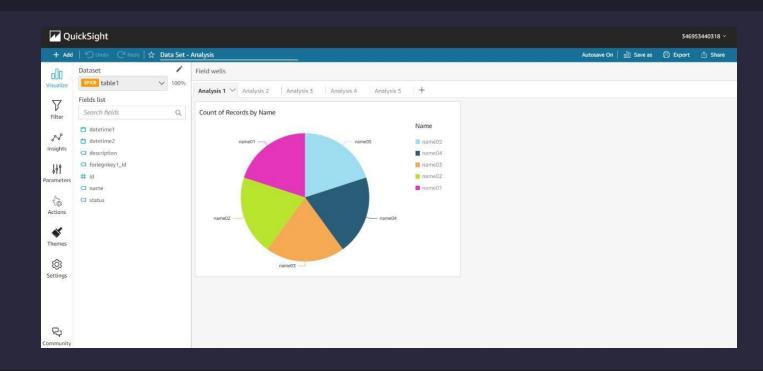






## QUICKSIGHT













## /CLOUDWATCH





#### /LOGS

 Collects compute performance metrics such as CPU, memory, and network as performance events



#### /METRICS

 CloudWatch Metric Streams enables you to create continuous, near-real-time streams of metrics to a destination of your choice



#### /MONITOR

 Create reusable graphs and visualize our cloud resources and applications in a unified view



#### /INSIGHTS

 Provides automated setup of observability for enterprise applications to get visibility into their health

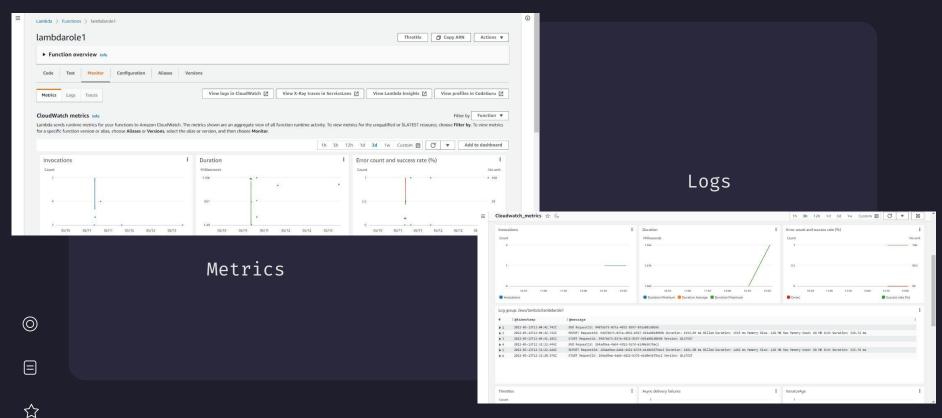








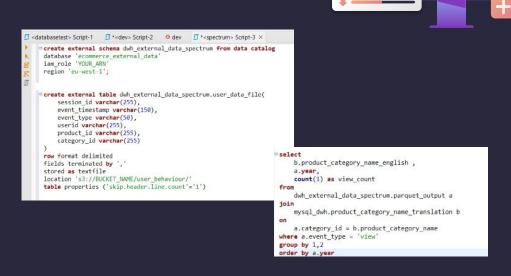
## **CLOUDWATCH**





## **REDSHIFT SPECTRUM**

Users can query data without having to load it.Multiple users can access same S3 bucket at same time.







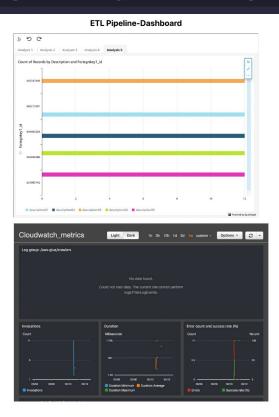


#### /Front-end

- 1. Build using ReactJS
- 2. Deployed in Netlify
- 3. Configured to a new domain
- 4. Added the domain in Domain and embedding section of the AWS page



## **CURRENTLY WORKING APPLICATION**











## **OPTIMIZATION TECHNIQUES FOR DATA**



#### **Compound sort keys**

Most frequently used columns put at front for speed



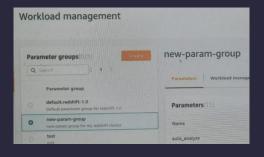
#### **Redshift DIST keys**

Data distribution across slices by Leader node matching values of a specific column



#### **Redshift Vacuuming**

Reclaims disk space occupied by rows marked for deletion in previous operations⇒ compacts table











## /INNOVATION

- Using Cloud based service -Amazon Web Services (AWS) for efficiency of the ETL Pipeline.
- AWS optimizes the software in terms of time, space and complexity.

Implemented using AWS Free Tier.







## /VIABILITY - INITIAL

<0.114s>

> TIME IN QUEUE

<0.627s>

> RUN TIME

<2.52KB>

> DATA SCANNED









## /VIABILITY - FINAL

<0.108s>

> TIME IN QUEUE

<0.446s>

> RUN TIME

<2.52KB>

> DATA SCANNED











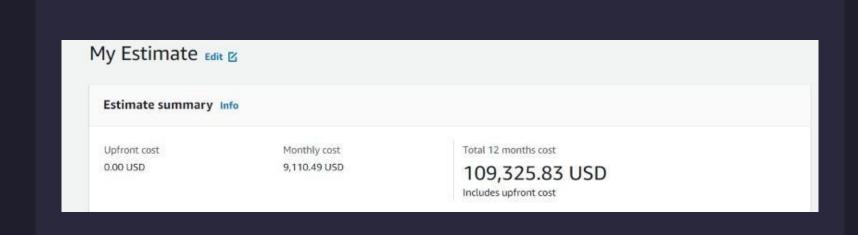


## **HOW REALISTIC IN THE MARKET?**





## /ROUGH ESTIMATE (\$USD)









## **/ROUGH ESTIMATE - SERVICES USED**

Service Name	•	Upfront cost	▽	Monthly cost	∇	Description	Region	⊽	Config Summary	▽
AWS Data Pipeline	ß	0.00 USD		117.00 USD		Data Pipeline	EU (Ireland)		Number of high	
AWS Data Transfer	ß	0.00 USD		40.96 USD		Data Transfer	EU (Ireland)		DT Inbound: All	
AWS Glue		0.00 USD		3.25 USD		Glue	EU (Ireland)		Number of DPU	
AWS Key Management Service	ß	0,00 USD		11.00 USD		KMS	EU (Ireland)		Number of cust	
AWS Lambda	ß	0.00 USD		0.00 USD		Lambda	EU (Ireland)		Architecture (x8	
Amazon Athena	ß	0.00 USD		74.27 USD		Athena	EU (Ireland)		Total number of	
Amazon CloudWatch	ß	0.00 USD		1.50 USD		Cloudwatch	EU (Ireland)		Number of Metri	
Amazon EC2	区	0.00 USD		71.12 USD		EC2	EU (Ireland)		Operating syste	
Amazon QuickSight	ß	0.00 USD		1,618.00 USD		Quicksight	EU (Ireland)		Number of work	
Amazon RDS Custom for SQL Server	ß	0,00 USD		2,412.36 USD		RDS	EU (Ireland)		Storage for each	
Amazon Redshift		0.00 USD		4,094.81 USD		Redshift	EU (Ireland)		Nodes ( 1 instan	
Amazon Simple Notification Service (SNS)	ß	0.00 USD		0.00 USD		SNS	EU (Ireland)			
Amazon Simple Storage Service (S3)	ß	0.00 USD		21.76 USD		S3	EU (Ireland)		S3 Standard sto	
Amazon Virtual Private Cloud (VPC)	ß	0.00 USD		644.46 USD		VPC	EU (Ireland)		Working days pe	







## **FUTURE PLANS**

#### **The Exterminators**

**ETL Pipelining + Dashboard + Error Logging** 



CSV File SQL Query

- •
- .







## **/FUTURE PLANS**

#### **The Exterminators**

**ETL Pipelining + Dashboard + Error Logging** 











## **/FUTURE PLANS**

#### The Exterminators

#### **ETL Pipelining + Dashboard + Error Logging**















